

SHIVEN RASTOGI

rastogishiven7@gmail.com | +91-7379889236 | github.com/Shivenrastogi59 | linkedin.com/in/shiven-rastogi

PROFESSIONAL SUMMARY

Computer Science graduate (AI & ML, CGPA 8.51) with strong foundations in data structures, algorithms, and system design. Experienced building high-throughput, low-latency distributed systems sustaining 50,000+ transactions/day and 10,000+ concurrent connections at sub-100ms P99 latency. Proficient in Python, Java, and JavaScript; skilled at profiling bottlenecks, designing event-driven microservices, and applying ML to real-world engineering problems at scale.

WORK EXPERIENCE

AI Developer, Nile Technologies Pvt. Ltd.

June 2026 - Present

- Architected a Warehouse Management SaaS processing 50,000+ inventory transactions/day; idempotent event-driven reconciliation engine cut manual operational overhead by 60% and held 99.9% uptime under peak concurrent load.
- Built a high-throughput logistics billing platform serving 200+ partner companies, processing 5,000+ billing events/day through a rate-card engine with queue-based job management and idempotent retries, ensuring zero data loss under concurrent load.
- Designed and optimised REST APIs and WebSocket feeds sustaining 10,000+ concurrent connections at sub-100ms P99 response times; profiled and eliminated bottlenecks across query plans, Redis caching, and async I/O paths.
- Deployed multi-agent AI orchestration (LangGraph, CrewAI, MCP) bridging LLM agents with live business data streams; reduced decision-pipeline latency by 45% and cut human intervention by 70% across production workflows.

AI Developer Intern, Nile Technologies Pvt. Ltd.

Oct 2025 - May 2026

- Built AIOps conversational systems (RAG, GraphRAG, Agentic RAG) on GCP and AWS using LangChain and CrewAI; reduced issue resolution time by 35% and cloud cost by 25% via model quantisation, request batching, and Redis caching.
- Optimised LLM inference pipelines behind load balancers; cut average response latency by 30% through async batching and tiered caching, directly improving SLA compliance for enterprise clients.

AI Research Intern, Yantra Pragya

May 2025 - Sep 2025

- Fine-tuned a causal language model on medical pathology corpora using domain-specific RAG and RLHF-style preference tuning; reduced hallucination rate by 30% on clinical QA benchmarks.
- Integrated the model into a real-time MERN interface with sub-500ms query response and stateful validation workflows for clinical report analysis.

PROJECTS

Hackathon Evaluation Automation *Python, Gemini 2.5 Pro, GPT-4o, LangChain, MERN*

- Led a 7-member team building a multi-LLM consensus engine with weighted scoring matrices and judge-bias correction; deployed at GLA University's Internal SIH to evaluate 90 teams across 6 judges simultaneously, cutting evaluation time by 80%.
- Designed a fault-tolerant pipeline with automatic retries and score normalisation to handle inconsistent LLM responses under load, maintaining evaluation integrity across 540+ individual team-judge pairs.

Real-Time Music Therapy Platform *YOLOv8, VGG16, Flask, MongoDB, React, Node.js*

- Built a 30 FPS emotion-detection pipeline (YOLOv8 + VGG16 ensemble) triggering generative music synthesis at sub-200ms end-to-end latency, with per-session emotional context tracking across concurrent real-time users.

EDUCATION

B.Tech - Computer Science & Engineering (AI & ML)

May 2026

GLA University, Mathura | CGPA: 8.51 / 10.0

TECHNICAL SKILLS

Languages: Python, Java, JavaScript, C, SQL, NoSQL

Systems & Infrastructure: Linux, Docker, AWS, GCP, Redis, Cloudflare, Microservices, WebSockets, REST APIs

CS Fundamentals: Data Structures & Algorithms, OOP, System Design, OS Internals, Computer Networks, DBMS

AI / ML: LangChain, LangGraph, CrewAI, RAG, Model Fine-tuning, Deep Learning, Neural Networks, MCP

ACHIEVEMENTS & RECOGNITION

- Coordinator, IBM ICE CodeCrunch National-Level Hackathon - managed end-to-end logistics and technical operations for 300+ participants across 12 colleges.
- Mentor (Certificate of Appreciation), Internal Smart India Hackathon 2025, GLAU - guided 6 teams through system architecture, code reviews, and live debugging under competition pressure.
- 1st Prize, Poetry Competition 'LEKHAK' - reflects precision in expression and strong written communication.